

ML/DL for Everyone Season2

with  TensorFlow

Lab 05-2 Logistic Regression cost function & optimizer

Code: <https://github.com/deeplearningzerotoall/TensorFlow>

Slides: <http://bit.ly/2LQMKvk>

Lecturer: SuSang Kim (healeess@kaist.ac.kr)




Logistic Regression

- What is Logistic Regression?
 - Classification
 - Logistic vs Linear
- How to solve?
 - Hypothesis Representation
 - Sigmoid/Logistic Function
 - Decision Boundary
 - Cost Function
 - Optimizer (Gradient Descent)
- Codes (Eager Execution)
- Summary

Cost Function

the cost function to fit the parameters(θ)



Given the training set how to we chose/fit θ ? $h_{\theta}(x) = y$ then Cost = 0




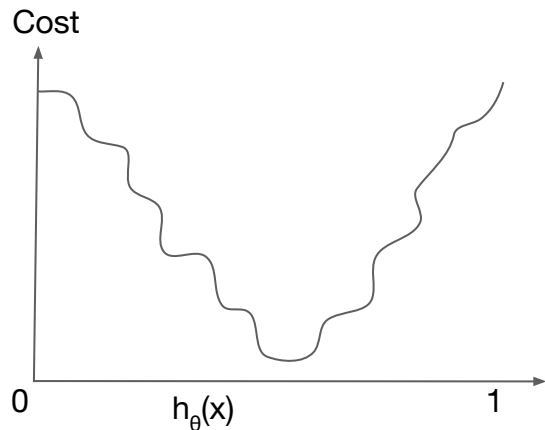
$$\text{cost}(h_{\theta}, (x), y) = -y \log(h_{\theta}(x)) - (1-y) \log(1 - h_{\theta}(x))$$

[Tensorflow Code]

```
def loss_fn(hypothesis, labels):  
    cost = -tf.reduce_mean(labels * tf.log(hypothesis) + (1 - labels) * tf.log(1 - hypothesis))  
    return cost
```

Cost Function

A convex logistic regression cost function



$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$


$$\text{cost}(h_{\theta}(x), y) = -y \log(h_{\theta}(x)) - (1-y) \log(1 - h_{\theta}(x))$$

[Tensorflow Code]

```
cost = -tf.reduce_mean(labels * tf.log(hypothesis) + (1 - labels) * tf.log(1 - hypothesis))
```

Optimization

How to minimize the cost function




$$\text{cost}(h_{\theta}(x), y) = -y \log(h_{\theta}(x)) - (1-y) \log(1 - h_{\theta}(x))$$

$$\text{Repeat } \left\{ \theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \right\}$$

[Tensorflow Code]

```
def grad(hypothesis, labels):  
    with tf.GradientTape() as tape:  
        loss_value = loss_fn(hypothesis, labels)  
    return tape.gradient(loss_value, [W,b])  
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.01)  
optimizer.apply_gradients(grads_and_vars=zip(grads,[W,b]))
```



Code(Eager)


```
import tensorflow.contrib.eager as tfe
tf.enable_eager_execution()
dataset = tf.data.Dataset.from_tensor_slices((x_train, y_train)).batch(len(x_train))
W = tf.Variable(tf.zeros([2,1]), name='weight')
b = tf.Variable(tf.zeros([1]), name='bias')
def logistic_regression(features):
    hypothesis = tf.div(1., 1. + tf.exp(tf.matmul(features, W) + b))
    return hypothesis
def loss_fn(hypothesis, labels):
    cost = -tf.reduce_mean(labels * tf.log(loss_fn(hypothesis)) + (1 - labels) * tf.log(1 - hypothesis))
    return cost
def grad(hypothesis, features, labels):
    with tf.GradientTape() as tape:
        loss_value = loss_fn(hypothesis, labels)
    return tape.gradient(loss_value, [W,b])
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.01)
for step in range(EPOCHS):
    for features, labels in tfe.Iterator(dataset):
        grads = grad(logistic_regression(features), features, labels)
        optimizer.apply_gradients(grads_and_vars=zip(grads,[W,b]))
        if step % 100 == 0:
            print("Iter: {}, Loss: {:.4f}".format(step, loss_fn(logistic_regression(features), labels)))

def accuracy_fn(hypothesis, labels):
    predicted = tf.cast(hypothesis > 0.5, dtype=tf.float32)
    accuracy = tf.reduce_mean(tf.cast(tf.equal(predicted, labels), dtype=tf.int32))
    return accuracy
test_acc = accuracy_fn(logistic_regression(x_test), y_test)
```





```
x_train = [[1., 2.],
            [2., 3.],
            [3., 1.],
            [4., 3.],
            [5., 3.],
            [6., 2.]]
y_train = [[0.],
            [0.],
            [0.],
            [1.],
            [1.],
            [1.]]

x_test = [[5., 2.]]
y_test = [[1.]]
```

Summary



$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$



What's Next?

Softmax classification

